

Staying up to Date with Online Content Changes Using Reinforcement Learning for Scheduling

Andrey Kolobov

Microsoft Research
akolobov@microsoft.com

Yuval Peres

yperes@gmail.com

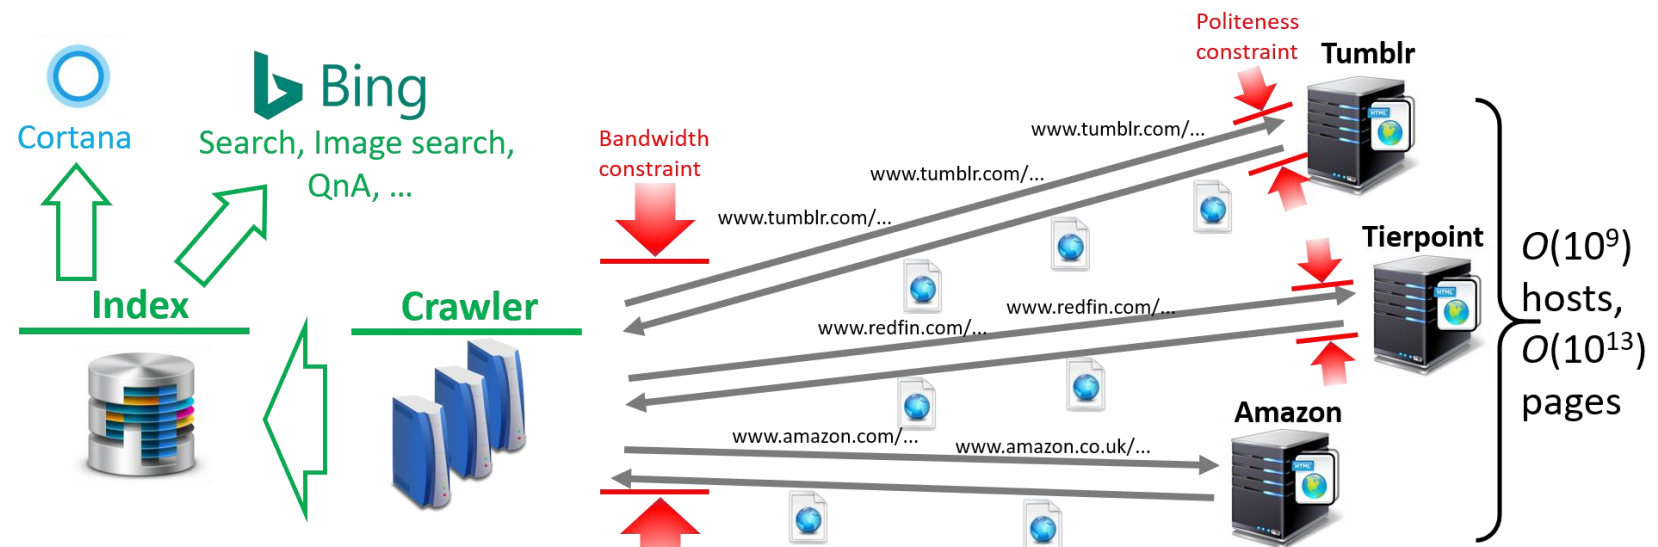
Cheng Lu

Microsoft Bing
cheng.lu@microsoft.com

Eric Horvitz

Microsoft Research
horvitz@microsoft.com

Web crawl scheduling



Fresh knowledge of online content is critical for search engines & other trackers.

They collect Web data by downloading (**crawling**) pages. Freshness crawl revisits indexed URLs to pick up changes.

How do we efficiently schedule crawl to maximize freshness across billions of pages under crawl bandwidth constraints?

Model

Set of sources W is fixed. Each source (e.g., URL) w in W has **importance score** μ_w and **Poisson change rate** Δ_w .

Want: policy π saying when to crawl each w in W to **minimize average importance-weighted staleness penalty**

$$J^\pi = \lim_{T \rightarrow \infty} \frac{1}{T} \mathbb{E}_{CrSeq \sim \pi, ChSeq \sim P(\tilde{\Delta})} \left[\frac{1}{T} \int_0^T \sum_{w \in W} \mu_w C(N_w(t)) dt \right]$$

Average over a long time period
Expectation over possible change and crawl sequences
 $C(N)$ is a staleness penalty function
 $N_w(t)$ is the number of times page w changed by time t since its latest crawl
subject to crawl bandwidth constraint R .

A popular $C(N)$ is **binary penalty** [3]: $C(0) = 0, C(N > 0) = 1$

Our Contributions


- New intuitive & efficiently optimizable objective for freshness crawling, *harmonic staleness penalty*.
- *LambdaCrawl* – an optimal efficient crawl scheduling algorithm for partial, full & mixed page change observability.
- *LambdaLearnAndCrawl* – RL for minimizing staleness under unknown page change rates Δ_w
- Approximation algorithm with data-dependent perf. guarantees

New Objective: Minimizing Harmonic Penalty

Let $C(N) = H_N = 1 + \frac{1}{2} + \dots + \frac{1}{N}$ if $N > 0$, else $C(0) = 0$.

H_N is the N -th harmonic number.

Under harmonic $C(N)$, J^π is equivalent to



Why specifically harmonic penalty?

- Has diminishing penalty property – makes intuitive sense.
- Guarantees that every source with change rate $\Delta_w > 0$ gets crawled occasionally (no such guarantee for binary $C(N)$!).
- Induces near-optimal policies w.r.t. binary $C(N)$ as well
- Efficient to optimize (unlike natural alternatives, e.g. $\log(1+N)$)

LambdaCrawl: Crawl Optimization Under Known Parameters

Incomplete Change Observations

For sources in set W^- , tracker finds out that they changed only when crawling them. Consider policies that crawl each w in W^- with some Poisson rate ρ_w .

Want: Crawl rates $\vec{\rho} = (\rho_w)_{w \in W^-}$ maximizing

$$\bar{J}^\pi = -J^\pi = \sum_{w \in W^-} \mu_w \ln \left(\frac{\rho_w}{\Delta_w + \rho_w} \right)$$

subject to

$$\sum_{w \in W^-} \rho_w = R, \rho_w \geq 0 \text{ for all } w \in W^-.$$

**Has a unique optimum for all $\rho_w \geq 0$.
Solve optimally using Lagrange multiplier method
+ binary search on multiplier λ**

Complete Change Observations

For sources in set W^0 , tracker gets notified about every change (e.g., using telemetry). Consider policies that crawl source w with probability p_w on each notification.

Want: Crawl probabilities $\vec{p} = (p_w)_{w \in W^0}$ maximizing

$$\bar{J}^\pi = -J^\pi = \sum_{w \in W^0} \mu_w \ln(p_w)$$

subject to

$$\sum_{w \in W^0} p_w \Delta_w = R, 0 \leq p_w \leq 1 \text{ for all } w \in W^0$$

**Has a unique optimum for all p_w in $[0,1]$.
Solve optimally by repeatedly applying Lagrange multiplier method to $O(|W^0|)$ inequality-constraint-free relaxations
(a bit tricky!)**

Mixed Observations

In reality, tracker must deal with a mix $W = W^- \cup W^0$
Want: Crawl rates $\vec{\rho} = (\rho_w)_{w \in W^-}$ and crawl probabilities $\vec{p} = (p_w)_{w \in W^0}$ maximizing

$$\bar{J}^\pi = -J^\pi = \sum_{w \in W^-} \mu_w \ln \left(\frac{\rho_w}{\Delta_w + \rho_w} \right) + \sum_{w \in W^0} \mu_w \ln(p_w)$$

subject to

$$\sum_{w \in W^-} \rho_w + \sum_{w \in W^0} p_w \Delta_w = R,$$

$\rho_w > 0$ for all $w \in W^-$, $0 < p_w \leq 1$ for all $w \in W^0$
**Boils down to finding the (unique!) optimal bandwidth split $R = R^- + R^0$
Easy concave maximization on $[0, R]$**

LambdaLearnAndCrawl: RL for Crawling

In reality, initially we don't know source change rates Δ_w !

Initialize Δ_w s to strictly positive values and...

Optimize policy π w.r.t. Δ_w estimates w/ LambdaCrawl.

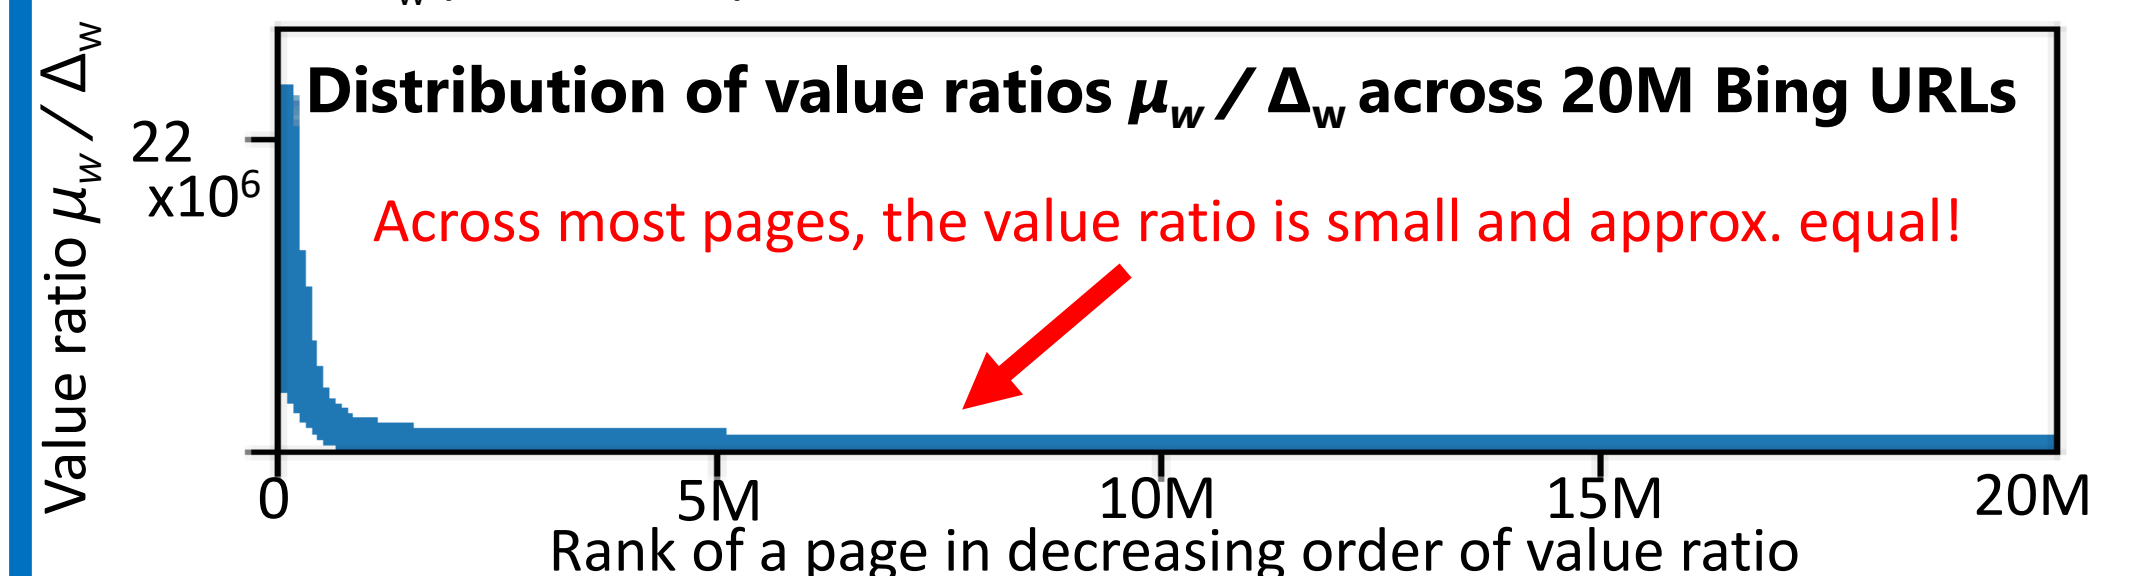
Run π for a few days.

Reestimate Δ_w s from new crawl & notification data using consistent estimators (see [2])

Guaranteed to converge to optimal π^* (assuming stationary Δ_w s).


Data-Dependent Approximation

One Δ_w parameter per source? That's a lot to learn! Can we avoid it?

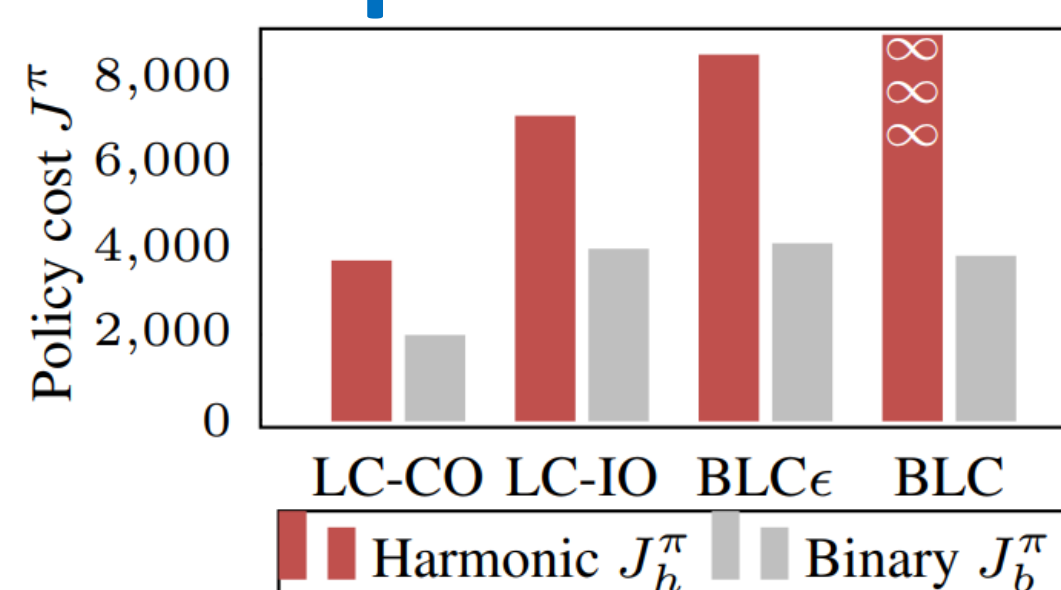


Theorem: across the constant value ratio region, the optimal crawl rate $\rho_w = \alpha \mu_w R$ – independent of change rate

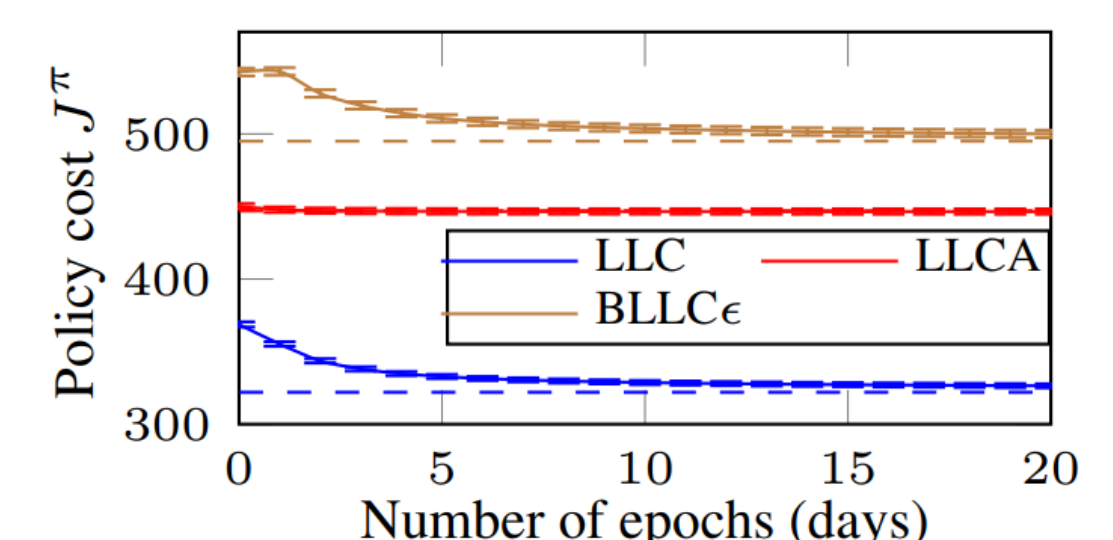
Experiments



LambdaCrawl (LC) outperforms BinaryLambdaCrawl (BLC) [1] w.r.t. harmonic penalty and even w.r.t. binary penalty, for which BLC is optimal under incomplete



Using complete observations as LambdaCrawl-CompObs (LC-CO) does yields huge performance gains



LambdaLearnAndCrawl converges quickly, and LambdaLearnAndCrawlApprox even faster (though to a worse solution)

References

- [1] Y. Azar, E. Horvitz, E. Lubetzky, Y. Peres, and D. Shahaf. "Tractable Near-optimal Policies for Crawling." PNAS-2018
- [2] J. Cho and H. Garcia-Molina. "Estimating frequency of change." ACM Transactions on Internet Technology, vol. 3(3), 2003
- [3] J. Cho and H. Garcia-Molina. "Effective page refresh policies for web crawlers." ACM Transactions on Database Systems, vol. 28 (4), 2003